

# ME766: Assignment 2

Mohd Safwan  
17D0700047

## Matrix Multiplication

- The simple  $\mathcal{O}(N^3)$  matrix multiplication algorithm is pretty straightforward to implement

```
void multiply(float A[N][N], float B[N][N], float C[N][N]) {  
    // Caculates C = AB where C is initialized to a zero matrix  
    int i, j, k;  
    for (i = 0; i < N; i++)  
        for (j = 0; j < N; j++)  
            for (k = 0; k < N; k++)  
                C[i][j] += A[i][k] * B[k][j];  
}
```

- Fixing the cache misses**

However we can see that `A[i][k]` is being used in the last loop but the value of `k` keeps changing on every iterations resulting in a **Cache Miss** everytime. Fixing this by swapping the last 2 loops gives us an almost **5x** speedup.

```
void multiply(float A[N][N], float B[N][N], float C[N][N]) {  
    // Caculates C = AB where C is initialized to a zero matrix  
    int i, j, k;  
    for (i = 0; i < N; i++)  
        for (k = 0; k < N; k++)  
            for (j = 0; j < N; j++)  
                C[i][j] += A[i][k] * B[k][j];  
}
```

- OMP Implementation**

Since there is no explicit data dependency, we can just distribute the operations of outer-most loop over threads arbitrarily.

```
void multiply(float A[N][N], float B[N][N], float C[N][N]) {  
    int i, j, k;  
    #pragma omp parallel for private(j, k)  
        for (i = 0; i < N; i++)  
            for (k = 0; k < N; k++)  
                for (j = 0; j < N; j++) C[i][j] += A[i][k] * B[k][j];  
}
```

- MPI Implementation**

- I have divided the  $A$  matrix into horizontal strips of almost equal lenghts, the number of division is equal to number of processors.
- The 0<sup>th</sup> node ( MASTER ) is responsible for doing the division and sending the strips of  $A$ . It also keeps one strip for itself.
- Each node/processor  $i$  receives a copy of  $B$  and the horizontal strip  $A_i$  of  $A$  that it has been assigned.
- Each node then calculates a strip of  $C$  that is  $C_i = A_i B$  and sends it back to MASTER .
- MASTER node collects all  $C_i$  strips and place them in the  $C$  matrix at correct locations.

## Gaussian Elimination

- Converting any invertible matrix to an upper triangular matrix is easy with the simple **Gauss Jordan Elimination** algorithm. The code is very straightforward for now.

```
void gaussianElimination(float R[N][N]) {  
    float r;  
    int i, j, k;  
    for (i = 0; i < N; i++) {  
        for (j = i + 1; j < N; j++) {  
            r = R[j][i] / R[i][i];  
            for (k = i; k < N; k++) R[j][k] -= r * R[i][k];  
        }  
    }  
}
```

- OMP Implementation**

There is obvious data dependency in the first loop, the elimination has to be done in correct order. The pivot must be chosen from first row, then second, and so on.

However, for all the rows below pivot row, we can parallelize the calculation of ratio and subtraction of rows.

```
void gaussianElimination(float R[N][N]) {  
    float r;  
    int i, j, k;  
    for (i = 0; i < N; i++) {  
        #pragma omp parallel for private(r, k)  
            for (j = i + 1; j < N; j++) {  
                r = R[j][i] / R[i][i];  
                for (k = i; k < N; k++) R[j][k] -= r * R[i][k];  
            }  
    }  
}
```

- MPI Implementation**

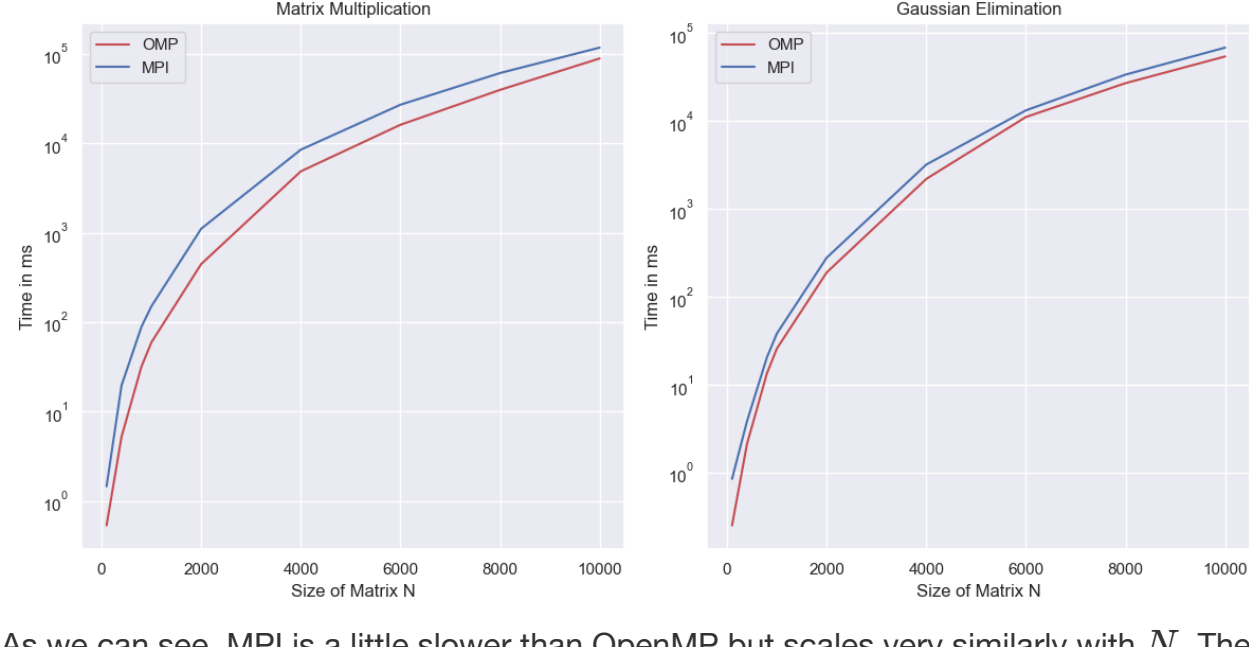
One obvious way to parallelize would be that when we want to subtract the pivot row from the rows below it, that operation can be distributed between nodes. However this would require division of the remaining rows  $\mathcal{O}(N)$  times and as many times sending the row strips to different nodes.

I have gone with a different approach.

- Every node is assigned a fixed set of rows from  $C$ , all the computations required for any of these rows will only be performed here.
- I keep track of how many rows to we need to modify in each node while doing the elimination.
- I have kept track of which node has the next pivot row, after each step of elimination, that row is asked by MASTER to send the next pivot row.
- MASTER receives the next pivot row from one of the nodes and places it in the  $C$  matrix at correct position.
- The work distribution is not uniform since more work is required to be done ton lower rows, but actually this method performs better since we only have to sed/receive  $\mathcal{O}(N)$  floats every iteration instead of  $\mathcal{O}(N^2)$  in the other method.


## Timing Analysis

- Scaling with  $N$**



As we can see, MPI is a little slower than OpenMP but scales very similarly with  $N$ . The largest problem size  $N = 10000$  was easily completed within 100 seconds.

- Scaling with Threads**



The plot for OMP Matrix Multiplication looks as expected, but there was a consistent jump at 5 threads while performing Gaussian Elimination. The same jump is present in both graphs for MPI as well. This may be due to the fact that after 4 physical cores, the framework uses hyperthreading for parallelization and thus there should be some overhead for that, which may be negligible for the first task.

## Hardware Information

- CPU**

```
Architecture:          x86_64  
CPU op-mode(s):        32-bit, 64-bit  
Byte Order:            Little Endian  
CPU(s):                8  
On-line CPU(s) list:   0-7  
Thread(s) per core:    2  
Core(s) per socket:    4  
Socket(s):             1  
NUMA node(s):          1  
Vendor ID:             GenuineIntel  
CPU family:            6  
Model:                 79  
Model name:            Intel(R) Xeon(R) CPU E5-2686 v4 @ 2.30GHz  
Stepping:              1  
CPU MHz:               2698.391  
CPU max MHz:           3000.0000  
CPU min MHz:           1200.0000  
BogoMIPS:              4600.18  
Hypervisor vendor:     Xen  
Virtualization type:   full  
L1d cache:             32K  
L1i cache:             32K  
L2 cache:              256K  
L3 cache:              46080K  
NUMA node0 CPU(s):     0-7
```

- OS**

```
NAME="Amazon Linux"  
VERSION="2"  
ID="amzn"  
ID_LIKE="centos rhel fedora"  
VERSION_ID="2"  
PRETTY_NAME="Amazon Linux 2"  
ANSI_COLOR="0;33"  
CPE_NAME="cpe:2.3:o:amazon:amazon_linux:2:-:internal"  
HOME_URL="https://amazonlinux.com/"  
VARIANT="internal"
```