

ME766: Assignment 3

Mohd Safwan
17D0700047

Implementation

The is the simple $\mathcal{O}(N^3)$ matrix multiplication algorithm.

```
void multiply(float A[N][N], float B[N][N], float C[N][N]) {
    // Caclulates C = AB where C is initialized to a zero matrix
    int i, j, k;
    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < N; k++)
                c[i][j] += A[i][k] * B[k][j];
}
```

For parallelizing it using CUDA, I have distributed the outer 2 for loops to parallel threads. Each thread computes the value C_{ij} or `c[i][j]` by performing a dot product.

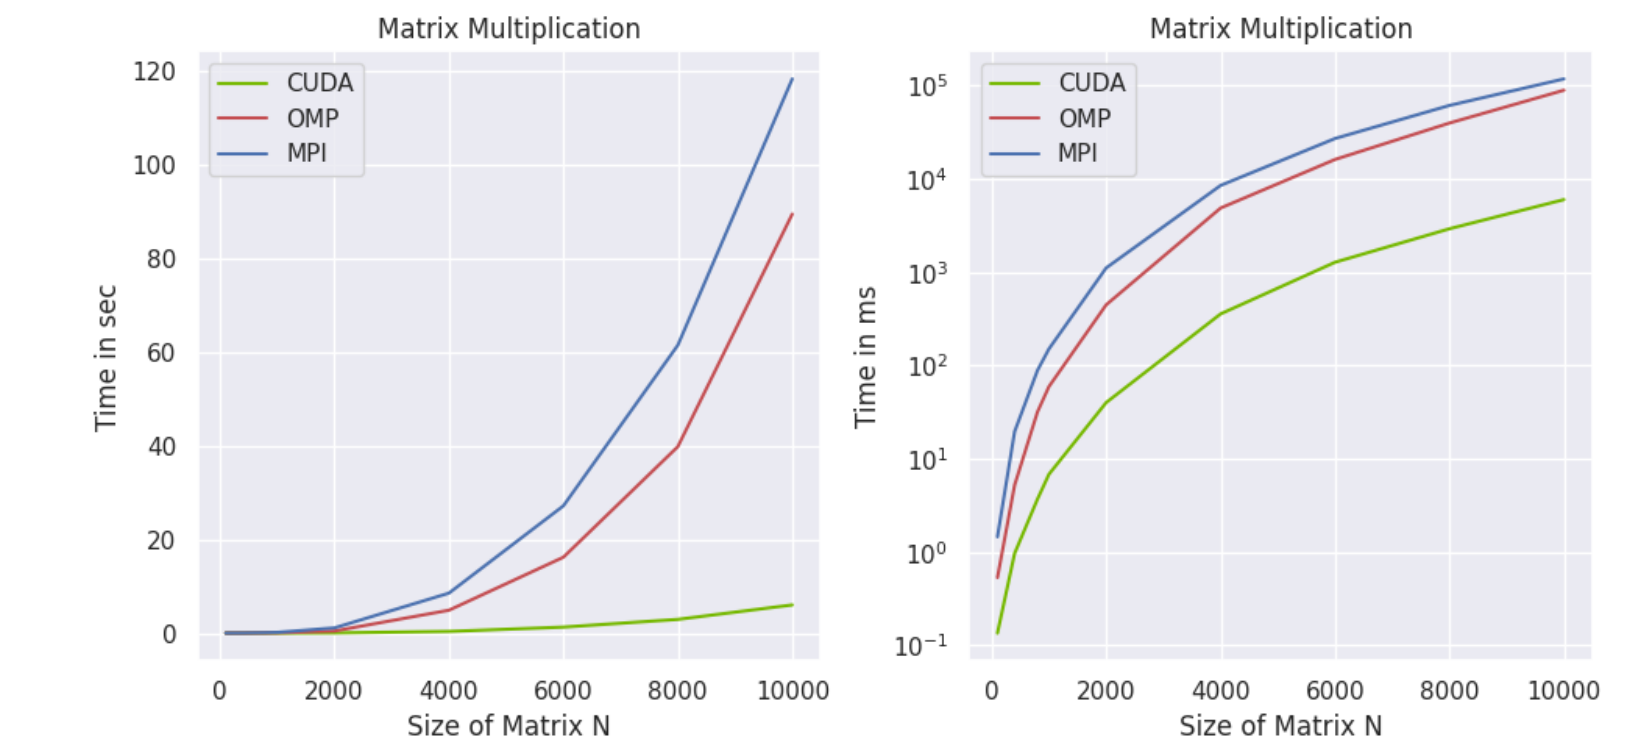
The CUDA implementation is pretty straightforward.

```
__global__ void multiply(float* A, float* B, float* C) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx >= N * N) return;
    int i = idx / N, j = idx % N;
    for (int k = 0; k < N; k++) C[i * N + j] += A[i * N + k] * B[k * N + j];
    return;
}
```

Timing Analysis

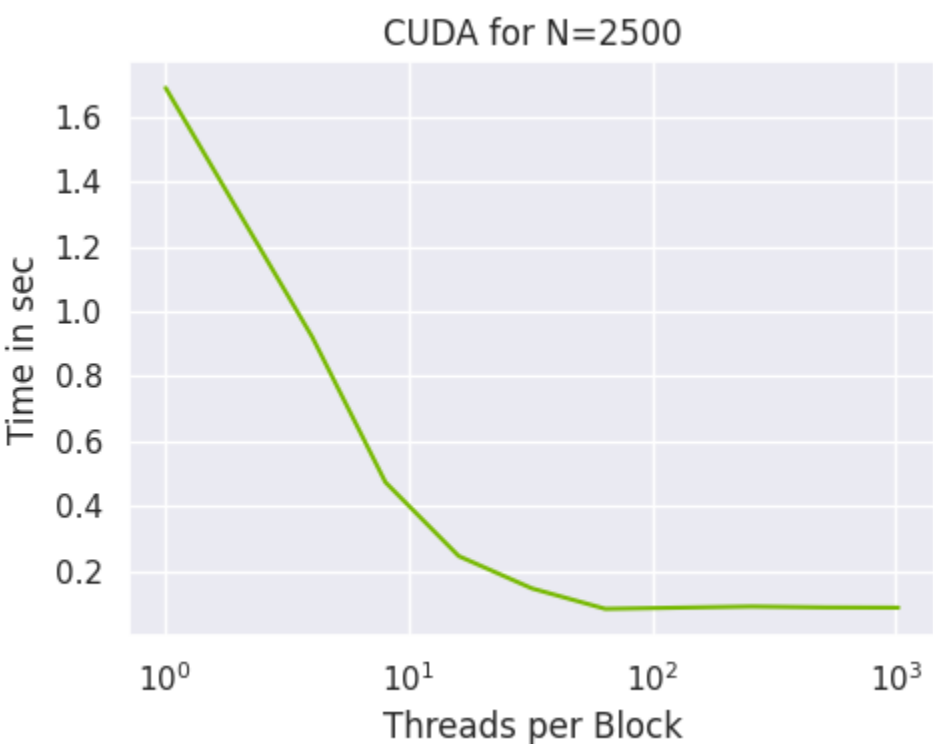
To be most fair in comparison, I have left out the time taken in initialising and allocating memory for A, B, C . The time raken in copying data between host and device and the matrix multiplication itself in only considered here

• Scaling with N



As we can see, CUDA scales very similarly with N and has the best performance as expected.

• Scaling with Threads per Block



The `TPB` parameter is very important as visible in the graph, above a certain threshold it gives more or less the same performance.

Hardware Information

CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "Tesla V100-SXM2-16GB"

CUDA Driver Version / Runtime Version

CUDA Capability Major/Minor version number:

Total amount of global memory:

(80) Multiprocessors, (64) CUDA Cores/MP:

GPU Max Clock rate:

Memory Clock rate:

Memory Bus Width:

L2 Cache Size:

Maximum Texture Dimension Size (x,y,z)

3D=(16384, 16384, 16384)

Maximum Layered 1D Texture Size, (num) layers

Maximum Layered 2D Texture Size, (num) layers

Total amount of constant memory:

Total amount of shared memory per block:

Total number of registers available per block:

Warp size:

Maximum number of threads per multiprocessor:

Maximum number of threads per block:

Max dimension size of a thread block (x,y,z):

Max dimension size of a grid size (x,y,z):

Maximum memory pitch:

Texture alignment:

Concurrent copy and kernel execution:

Run time limit on kernels:

Integrated GPU sharing Host Memory:

Support host page-locked memory mapping:

Alignment requirement for Surfaces:

Device has ECC support:

Device supports Unified Addressing (UVA):

Device supports Compute Preemption:

Supports Cooperative Kernel Launch:

Supports MultiDevice Co-op Kernel Launch:

Device PCI Domain ID / Bus ID / location ID:

Compute Mode:

10.1 / 10.0

7.0

16130 MBytes (16914055168 bytes)

5120 CUDA Cores

1530 MHz (1.53 GHz)

877 Mhz

4096-bit

6291456 bytes

1D=(131072), 2D=(131072, 65536),

1D=(32768), 2048 layers

2D=(32768, 32768), 2048 layers

65536 bytes

49152 bytes

65536

32

2048

1024

(1024, 1024, 64)

(2147483647, 65535, 65535)

2147483647 bytes

512 bytes

Yes with 2 copy engine(s)

No

No

Yes

Yes

Enabled

Yes

Yes

Yes

Yes

0 / 0 / 30

< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 10.1,
CUDA Runtime Version = 10.0, NumDevs = 1
Result = PASS